

News Popularity: Predicting popular news articles.

```
In [127... pwd
```

```
Out[127]: '/Users/pratik'
```

```
In [128... import pandas as pd
```

Collecting Data

```
In [129... news = pd.read_csv('OnlineNewsPopularity_for_python.csv')
```

```
In [130... news.head(4)
```

```
Out[130]:
```

	url	timedelta	n_tokens_title	n_tokens_content	n_unique_tokens
0	http://mashable.com/2013/01/07/amazon-instant-...	731	12	219	0.663594
1	http://mashable.com/2013/01/07/ap-samsung-spon...	731	9	255	0.604743
2	http://mashable.com/2013/01/07/apple-40-billio...	731	9	211	0.575130
3	http://mashable.com/2013/01/07/astronaut-notre...	731	9	531	0.503788

4 rows x 61 columns

```
In [131... # handle goal attrubte to binary
popular = news.shares >= 1400
unpopular = news.shares < 1400
news.loc[popular, 'shares'] = 1
news.loc[unpopular, 'shares'] = 0
```

```
In [132... news.head(4)
```

```
Out[132]:
```

	url	timedelta	n_tokens_title	n_tokens_content	n_unique_tokens
0	http://mashable.com/2013/01/07/amazon-instant-...	731	12	219	0.663594
1	http://mashable.com/2013/01/07/ap-samsung-spon...	731	9	255	0.604743
2	http://mashable.com/2013/01/07/apple-40-billio...	731	9	211	0.575130
3	http://mashable.com/2013/01/07/astronaut-notre...	731	9	531	0.503788

4 rows x 61 columns

Decision Tree

```
In [133... print(news.columns)
news.info()
```

```
Index(['url', 'timedelta', 'n_tokens_title', 'n_tokens_content',
      'n_unique_tokens', 'n_non_stop_words', 'n_non_stop_unique_tokens',
      'num_hrefs', 'num_self_hrefs', 'num_imgs', 'num_videos',
      'average_token_length', 'num_keywords', 'data_channel_is_lifestyle',
      'data_channel_is_entertainment', 'data_channel_is_bus',
      'data_channel_is_socmed', 'data_channel_is_tech',
      'data_channel_is_world', 'kw_min_min', 'kw_max_min', 'kw_avg_min',
      'kw_min_max', 'kw_max_max', 'kw_avg_max', 'kw_min_avg', 'kw_max_avg',
      'kw_avg_avg', 'self_reference_min_shares', 'self_reference_max_shares',
      'self_reference_avg_shares', 'weekday_is_monday', 'weekday_is_tuesday',
      'weekday_is_wednesday', 'weekday_is_thursday', 'weekday_is_friday',
      'weekday_is_saturday', 'weekday_is_sunday', 'is_weekend', 'LDA_00',
      'LDA_01', 'LDA_02', 'LDA_03', 'LDA_04', 'global_subjectivity',
      'global_sentiment_polarity', 'global_rate_positive_words',
      'global_rate_negative_words', 'rate_positive_words',
      'rate_negative_words', 'avg_positive_polarity', 'min_positive_polarity',
      'max_positive_polarity', 'avg_negative_polarity',
      'min_negative_polarity', 'max_negative_polarity', 'title_subjectivity',
      'title_sentiment_polarity', 'abs_title_subjectivity',
      'abs_title_sentiment_polarity', 'shares'],
      dtype='object')
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 39644 entries, 0 to 39643
```

```
Data columns (total 61 columns):
```

#	Column	Non-Null Count	Dtype
0	url	39644 non-null	object
1	timedelta	39644 non-null	int64
2	n_tokens_title	39644 non-null	int64
3	n_tokens_content	39644 non-null	int64
4	n_unique_tokens	39644 non-null	float64
5	n_non_stop_words	39644 non-null	float64
6	n_non_stop_unique_tokens	39644 non-null	float64
7	num_hrefs	39644 non-null	int64
8	num_self_hrefs	39644 non-null	int64
9	num_imgs	39644 non-null	int64
10	num_videos	39644 non-null	int64
11	average_token_length	39644 non-null	float64
12	num_keywords	39644 non-null	int64
13	data_channel_is_lifestyle	39644 non-null	int64
14	data_channel_is_entertainment	39644 non-null	int64
15	data_channel_is_bus	39644 non-null	int64
16	data_channel_is_socmed	39644 non-null	int64
17	data_channel_is_tech	39644 non-null	int64
18	data_channel_is_world	39644 non-null	int64
19	kw_min_min	39644 non-null	int64
20	kw_max_min	39644 non-null	float64
21	kw_avg_min	39644 non-null	float64
22	kw_min_max	39644 non-null	int64
23	kw_max_max	39644 non-null	int64
24	kw_avg_max	39644 non-null	float64
25	kw_min_avg	39644 non-null	float64
26	kw_max_avg	39644 non-null	float64
27	kw_avg_avg	39644 non-null	float64
28	self_reference_min_shares	39644 non-null	float64
29	self_reference_max_shares	39644 non-null	float64
30	self_reference_avg_shares	39644 non-null	float64
31	weekday_is_monday	39644 non-null	int64
32	weekday_is_tuesday	39644 non-null	int64
33	weekday_is_wednesday	39644 non-null	int64
34	weekday_is_thursday	39644 non-null	int64
35	weekday_is_friday	39644 non-null	int64
36	weekday_is_saturday	39644 non-null	int64
37	weekday_is_sunday	39644 non-null	int64
38	is_weekend	39644 non-null	int64
39	LDA_00	39644 non-null	float64

```

40 LDA_01 39644 non-null float64
41 LDA_02 39644 non-null float64
42 LDA_03 39644 non-null float64
43 LDA_04 39644 non-null float64
44 global_subjectivity 39644 non-null float64
45 global_sentiment_polarity 39644 non-null float64
46 global_rate_positive_words 39644 non-null float64
47 global_rate_negative_words 39644 non-null float64
48 rate_positive_words 39644 non-null float64
49 rate_negative_words 39644 non-null float64
50 avg_positive_polarity 39644 non-null float64
51 min_positive_polarity 39644 non-null float64
52 max_positive_polarity 39644 non-null float64
53 avg_negative_polarity 39644 non-null float64
54 min_negative_polarity 39644 non-null float64
55 max_negative_polarity 39644 non-null float64
56 title_subjectivity 39644 non-null float64
57 title_sentiment_polarity 39644 non-null float64
58 abs_title_subjectivity 39644 non-null float64
59 abs_title_sentiment_polarity 39644 non-null float64
60 shares 39644 non-null int64
dtypes: float64(34), int64(26), object(1)
memory usage: 18.5+ MB

```

Exploring and Preparing the Data

```
In [134... target = news['shares']
```

Train, Test data with a split of 75% data for training set and 25% data for testing set and seed of 23458. Dropping the "url" column as it is a string type object and just the address of the site, dropping "shares" column from X as "shares" is a response or target variable.

```
In [135... from sklearn.model_selection import train_test_split
y = target
x = news.drop(['shares', 'url'], axis=1)
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.25, random_state=23458)
```

Design Decision Tree

```
In [136... from sklearn import tree
from sklearn.tree import DecisionTreeClassifier
```

```
In [137... model = tree.DecisionTreeClassifier()
model = model.fit(x_train, y_train)
```

Evaluating the model

```
In [138... from sklearn.metrics import confusion_matrix
from sklearn.metrics import accuracy_score
y_predict = model.predict(x_test)
print(confusion_matrix(y_test, y_predict))
```

```
[[2513 2050]
 [2098 3250]]
```

```
In [139... print(accuracy_score(y_test, y_predict)*100)
```

```
58.147512864494
```

The results show the accuracy of the model to be 58% for seed of 23458.

Random Forest

```
In [140... y = target
x = news.drop(['shares', 'url'], axis=1)
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.25, random_state=2
```

```
In [141... from sklearn.ensemble import RandomForestClassifier
clf = RandomForestClassifier()
clf.fit(x_train, y_train)
```

```
Out[141]: RandomForestClassifier()
```

```
In [142... y_predict = clf.predict(x_test)
print(confusion_matrix(y_test, y_predict))
```

```
[[2794 1769]
 [1520 3828]]
```

```
In [143... print(accuracy_score(y_test, y_predict)*100)
```

```
66.81465038845728
```

The Results show that using Random Forest the accuracy of predicting popular and unpoularnews based on the model is 67% based on ratio of (True positives and True Negatives) to all predicted observations (TP+FN+FP+TN)

Important Features

```
In [144... import pandas as pd
feature_importances = pd.DataFrame(clf.feature_importances_, index = news.columns[1:60],
feature_importances
```

```
Out[144]:
```

	importance
kw_avg_avg	0.040599
kw_max_avg	0.038371
LDA_02	0.033017
timedelta	0.031161
self_reference_avg_sharess	0.029585
LDA_01	0.028522
LDA_04	0.028398
self_reference_min_shares	0.027897
kw_avg_min	0.027684
LDA_00	0.027673
kw_avg_max	0.027526
global_subjectivity	0.026878
n_unique_tokens	0.026258
n_non_stop_unique_tokens	0.026078
average_token_length	0.025736
kw_max_min	0.025464

n_tokens_content	0.025322
LDA_03	0.025283
global_rate_positive_words	0.024946
avg_positive_polarity	0.024771
kw_min_avg	0.024475
global_sentiment_polarity	0.023800
self_reference_max_shares	0.023728
avg_negative_polarity	0.022924
global_rate_negative_words	0.022820
num_hrefs	0.020784
rate_negative_words	0.019936
rate_positive_words	0.019703
kw_min_max	0.016427
n_tokens_title	0.015610
title_sentiment_polarity	0.014466
min_negative_polarity	0.014348
num_imgs	0.013600
max_negative_polarity	0.013421
num_self_hrefs	0.012824
min_positive_polarity	0.012670
title_subjectivity	0.012334
abs_title_subjectivity	0.011780
abs_title_sentiment_polarity	0.011593
is_weekend	0.011208
max_positive_polarity	0.010282
num_keywords	0.010244
data_channel_is_entertainment	0.009047
num_videos	0.007659
data_channel_is_world	0.006367
kw_max_max	0.005807
data_channel_is_socmed	0.005328
data_channel_is_tech	0.004845
kw_min_min	0.004643
weekday_is_saturday	0.004358
weekday_is_tuesday	0.003403
weekday_is_thursday	0.002950
weekday_is_wednesday	0.002925
weekday_is_friday	0.002870
weekday_is_monday	0.002835

weekday_is_sunday	0.002627
data_channel_is_bus	0.002405
data_channel_is_lifestyle	0.001652
n_non_stop_words	0.000134

We see that the top five most important features for the model to predict popular and unpopular news are kw_avg_avg, kw_max_avg, timedelta, LDA_02 and self_reference_min_shares in that order.

Finding Correlation

```
In [145]: import numpy as np
np.corrcoef(y_test, y_predict)
```

```
Out[145]: array([[1.          , 0.3298427],
                [0.3298427, 1.          ]])
```

We see that the correlation between actual and predicted values is 33%

Finding RMSE

```
In [146]: from sklearn.metrics import mean_squared_error
rmse = mean_squared_error(y_test, y_predict, squared=False)
```

```
In [147]: print(rmse)
```

```
0.5760672670057094
```

The results shows that the root mean squared error is 0.57. It appears that the model is a little less effective than the acceptable range as the difference between the actual and predicted test scores or the errors, account for more than 0.50 or 50%. The lower the rmse value the better the model can be assumed to fit the dataset (rmse values lie in range from 0 to 1).

Part II: Dimensionality Reduction with correlation graph


```
In [148]: target = news["shares"]
y = target
x = news.drop(['shares', 'url'], axis = 1)
```

Correlation Graph

```
In [149]: import numpy as np
import matplotlib.pyplot as plt

corr = x.corr()
fig = plt.figure()
ax = fig.add_subplot(111)
cax = ax.matshow(corr, cmap='coolwarm', vmin=-1, vmax=1)
fig.colorbar(cax)
ticks = np.arange(0, len(x.columns), 1)
ax.set_xticks(ticks)
plt.xticks(rotation=90)
ax.set_yticks(ticks)
ax.set_xticklabels(x.columns)
```

```
ax.set_yticklabels(x.columns)
plt.show()
```



```
In [150]: corr_matrix = x.corr().abs()
corr_matrix
```

```
Out[150]:
```

	timedelta	n_tokens_title	n_tokens_content	n_unique_tokens	n_non_stop
timedelta	1.000000	0.240320	0.062867	0.002866	0.
n_tokens_title	0.240320	1.000000	0.018160	0.005318	0
n_tokens_content	0.062867	0.018160	1.000000	0.004737	(
n_unique_tokens	0.002866	0.005318	0.004737	1.000000	0
n_non_stop_words	0.000089	0.004754	0.017512	0.999572	1.
n_non_stop_unique_tokens	0.003805	0.005420	0.000373	0.999852	0
num_hrefs	0.000832	0.053496	0.423065	0.004352	0
num_self_hrefs	0.064530	0.014856	0.304682	0.006620	0
num_imgs	0.027636	0.008858	0.342600	0.018802	0.
num_videos	0.000936	0.051460	0.103699	0.000597	0.
average_token_length	0.130465	0.071403	0.167789	0.026407	0

	num_keywords	0.046884	0.006077	0.072845	0.003679	0
	data_channel_is_lifestyle	0.054492	0.070815	0.037548	0.001653	0
	data_channel_is_entertainment	0.049109	0.132791	0.060200	0.011016	0
	data_channel_is_bus	0.055788	0.023902	0.006105	0.000264	0
	data_channel_is_socmed	0.076287	0.090394	0.033424	0.000945	0
	data_channel_is_tech	0.083277	0.046716	0.025408	0.002328	0
	data_channel_is_world	0.170250	0.049223	0.055989	0.005535	0
	kw_min_min	0.591199	0.110672	0.054345	0.001601	0
	kw_max_min	0.029503	0.005890	0.000066	0.000552	0
	kw_avg_min	0.133225	0.031400	0.003545	0.000826	0
	kw_min_max	0.076590	0.012926	0.022786	0.000577	0
	kw_max_max	0.637824	0.120841	0.058860	0.001624	0
	kw_avg_max	0.493093	0.115746	0.096460	0.000805	0
	kw_min_avg	0.157204	0.002370	0.022286	0.004563	0
	kw_max_avg	0.051820	0.006918	0.030496	0.002120	0
	kw_avg_avg	0.163164	0.004296	0.079624	0.002083	0
	self_reference_min_shares	0.011438	0.004563	0.030686	0.001036	0
	self_reference_max_shares	0.014501	0.000128	0.025657	0.000222	0
	self_reference_avg_sharess	0.015655	0.000661	0.013809	0.001992	0
	weekday_is_monday	0.006129	0.004274	0.002484	0.002142	0
	weekday_is_tuesday	0.005781	0.009322	0.004027	0.010538	0
	weekday_is_wednesday	0.009961	0.008935	0.016891	0.002224	0
	weekday_is_thursday	0.004042	0.015472	0.007395	0.002248	0
	weekday_is_friday	0.002853	0.002015	0.015949	0.001398	0
	weekday_is_saturday	0.004067	0.015013	0.034538	0.002563	0
	weekday_is_sunday	0.004226	0.006289	0.036394	0.001803	0
	is_weekend	0.000272	0.005996	0.052024	0.003186	0
	LDA_00	0.080894	0.070038	0.026218	0.002213	0
	LDA_01	0.004423	0.063568	0.009724	0.000827	0
	LDA_02	0.141713	0.038365	0.087266	0.006855	0
	LDA_03	0.030838	0.042208	0.140141	0.003689	0
	LDA_04	0.092906	0.065063	0.041265	0.004260	0
	global_subjectivity	0.133837	0.056804	0.127879	0.000180	0
	global_sentiment_polarity	0.158646	0.072226	0.021937	0.000523	0
	global_rate_positive_words	0.207604	0.064951	0.133979	0.000014	0
	global_rate_negative_words	0.010266	0.015530	0.125013	0.000877	0
	rate_positive_words	0.198654	0.066589	0.098960	0.000667	0
	rate_negative_words	0.071968	0.034186	0.101053	0.001657	0
	avg_positive_polarity	0.126344	0.049619	0.135123	0.000487	0

min_positive_polarity	0.054772	0.025069	0.261493	0.009193	0.
max_positive_polarity	0.098288	0.021662	0.415706	0.009054	0
avg_negative_polarity	0.000507	0.017096	0.130375	0.001453	C
min_negative_polarity	0.062175	0.029146	0.450603	0.009902	C
max_negative_polarity	0.063239	0.011425	0.225870	0.007315	C
title_subjectivity	0.015919	0.077245	0.004484	0.004678	0.
title_sentiment_polarity	0.038711	0.000240	0.023358	0.002333	0
abs_title_subjectivity	0.011551	0.146954	0.007136	0.009242	C
abs_title_sentiment_polarity	0.002745	0.040550	0.013439	0.004217	0

59 rows × 59 columns

Select Upper triangle of correlation matrix

```
In [151... import numpy as np
upper = corr_matrix.where(np.triu(np.ones(corr_matrix.shape) , k=1).astype(np.bool))

/var/folders/fp/szpx3gxd4ln849f9dbx_lqw00000gn/T/ipykernel_1107/2665990962.py:2: Deprecati
onWarning: `np.bool` is a deprecated alias for the builtin `bool`. To silence this warn
ing, use `bool` by itself. Doing this will not modify any behavior and is safe. If you
specifically wanted the numpy scalar type, use `np.bool_` here.
Deprecated in NumPy 1.20; for more details and guidance: https://numpy.org/devdocs/relea
se/1.20.0-notes.html#deprecations
upper = corr_matrix.where(np.triu(np.ones(corr_matrix.shape) , k=1).astype(np.bool))
```

Find index of feature columns with correlation greater than 0.95

```
In [152... to_drop = [column for column in upper.columns if any(upper[column] > 0.6)]
xnew = x.drop(to_drop, axis = 1)
xnew
```

Out[152]:

	timedelta	n_tokens_title	n_tokens_content	n_unique_tokens	num_hrefs	num_self_hrefs	num_i
0	731	12	219	0.663594	4	2	
1	731	9	255	0.604743	3	1	
2	731	9	211	0.575130	3	1	
3	731	9	531	0.503788	9	0	
4	731	13	1072	0.415646	19	19	
...
39639	8	11	346	0.529052	9	7	
39640	8	12	328	0.696296	9	7	
39641	8	10	442	0.516355	24	1	
39642	8	6	682	0.539493	10	1	
39643	8	10	157	0.701987	1	1	

39644 rows × 43 columns

New Accuracy with new Xnew(after dropping highly correlated variables)

Using Xnew(after dropping highly correlated var) for randomising, Splitting, Training model and evaluating

In [153]:

```
xnew
```

Out[153]:

	timedelta	n_tokens_title	n_tokens_content	n_unique_tokens	num_hrefs	num_self_hrefs	num_i
0	731	12	219	0.663594	4	2	
1	731	9	255	0.604743	3	1	
2	731	9	211	0.575130	3	1	
3	731	9	531	0.503788	9	0	
4	731	13	1072	0.415646	19	19	
...
39639	8	11	346	0.529052	9	7	
39640	8	12	328	0.696296	9	7	
39641	8	10	442	0.516355	24	1	
39642	8	6	682	0.539493	10	1	
39643	8	10	157	0.701987	1	1	

39644 rows x 43 columns

In [154]:

```
to_drop
```

Out[154]:

```
['n_non_stop_words',  
 'n_non_stop_unique_tokens',  
 'kw_avg_min',  
 'kw_max_max',  
 'kw_avg_avg',  
 'self_reference_avg_sharess',  
 'is_weekend',  
 'LDA_00',  
 'LDA_02',  
 'LDA_04',  
 'rate_positive_words',  
 'rate_negative_words',  
 'avg_positive_polarity',  
 'max_positive_polarity',  
 'min_negative_polarity',  
 'abs_title_sentiment_polarity']
```

Dropping to_drop columns

In [155]:

```
news_new = news.drop(['n_non_stop_words',  
 'n_non_stop_unique_tokens',  
 'kw_avg_min',  
 'kw_max_max',  
 'kw_avg_avg',  
 'self_reference_avg_sharess',  
 'is_weekend',  
 'LDA_00',  
 'LDA_02',  
 'LDA_04',  
 'rate_positive_words',  
 'rate_negative_words',  
 'avg_positive_polarity',  
 'max_positive_polarity',
```

```
'min_negative_polarity',
'abs_title_sentiment_polarity'], axis = 1)
```

As above, the `to_drop` function showed high correlation for 16 columns, which we will remove from our dataset for further analysis. We can see the remaining variables in `xnew` as below, which we will consider for further analysis

In [156... `xnew`

Out[156]:

	timedelta	n_tokens_title	n_tokens_content	n_unique_tokens	num_hrefs	num_self_hrefs	num_i
0	731	12	219	0.663594	4	2	
1	731	9	255	0.604743	3	1	
2	731	9	211	0.575130	3	1	
3	731	9	531	0.503788	9	0	
4	731	13	1072	0.415646	19	19	
...	
39639	8	11	346	0.529052	9	7	
39640	8	12	328	0.696296	9	7	
39641	8	10	442	0.516355	24	1	
39642	8	6	682	0.539493	10	1	
39643	8	10	157	0.701987	1	1	

39644 rows x 43 columns

In [157... `news_new`

Out[157]:

	url	timedelta	n_tokens_title	n_tokens_content	n_unique_tok
0	http://mashable.com/2013/01/07/amazon-instant-...	731	12	219	0.663
1	http://mashable.com/2013/01/07/ap-samsung-spon...	731	9	255	0.604
2	http://mashable.com/2013/01/07/apple-40-billio...	731	9	211	0.575
3	http://mashable.com/2013/01/07/astronaut-notre...	731	9	531	0.503
4	http://mashable.com/2013/01/07/att-u-verse-apps/	731	13	1072	0.415
...
39639	http://mashable.com/2014/12/27/samsung-app-aut...	8	11	346	0.529
39640	http://mashable.com/2014/12/27/seth-rogen-jame...	8	12	328	0.696
39641	http://mashable.com/2014/12/27/son-pays-off-mo...	8	10	442	0.516
39642	http://mashable.com/2014/12/27/ukraine-blasts/	8	6	682	0.539
39643	http://mashable.com/2014/12/27/youtube-channel...	8	10	157	0.701

Decision Tree Model: (After dropping the correlated variables)

(Note: It is not necessary to normalize or standardize the data for Decision Tree and Random Forest models)

```
In [158... y = target
x = xnew
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.25, random_state=2
```

```
In [159... from sklearn import tree
from sklearn.tree import DecisionTreeClassifier
model = tree.DecisionTreeClassifier()
model = model.fit(x_train, y_train)
y_predict = model.predict(x_test)
print(confusion_matrix(y_test, y_predict))

[[2503 2060]
 [2151 3197]]
```

```
In [160... print(accuracy_score(y_test, y_predict)*100)

57.51185551407527
```

Random Forest Model: (After dropping the correlated variables)

```
In [161... y = target
x = xnew
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.25, random_state=2
```

```
In [162... from sklearn.ensemble import RandomForestClassifier
model = RandomForestClassifier()
model = model.fit(x_train, y_train)
y_predict = model.predict(x_test)
print(confusion_matrix(y_test, y_predict))

[[2791 1772]
 [1509 3839]]
```

```
In [163... print(accuracy_score(y_test, y_predict)*100)

66.89536878216124
```

It is observed that there isn't much difference in the accuracies of Decision Tree and Random Forest models after dropping the correlated variables.

Normalizing the data(Normalization: Use MinMaxScaler from sklearn, Standardization: Use StandardScaler from sklearn)

```
In [164... from sklearn.preprocessing import MinMaxScaler
scaler = MinMaxScaler()
scaled = scaler.fit_transform(xnew)
print(scaled)

[[1.          0.47619048  0.02584376  ...  0.5          0.40625        0.          ]
 [1.          0.33333333  0.03009205  ...  0.          0.5          1.          ]
 [1.          0.33333333  0.02489969  ...  0.          0.5          1.          ]
 ...
 [0.          0.38095238  0.05215955  ...  0.45454545  0.56818182  0.09090909]]
```

```
[0.          0.19047619  0.08048147 ...  0.          0.5          1.          ]
[0.          0.38095238  0.01852726 ...  0.33333333  0.625         0.33333333]]
```

Training(90) and Test(10) Split: (Note : Use Normalized x (xnew) when splitting data)

```
In [165... y = target
x = scaled
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size = 0.10, random_state
y_train
```

```
Out[165]: 24329      1
11705      0
1995       1
22132      1
24562      1
..
32399      0
17048      1
23924      0
34086      0
27439      0
Name: shares, Length: 35679, dtype: int64
```

```
In [166... x_train.shape
```

```
Out[166]: (35679, 43)
```

```
In [167... x_test.shape
```

```
Out[167]: (3965, 43)
```

Train the model, Use Linear Separator to design the model

Designing the model

```
In [168... from sklearn import svm
clf = svm.SVC(kernel = "linear")
clf
```

```
Out[168]: SVC(kernel='linear')
```

```
In [169... clf.fit(x_train, y_train)
y_predict = clf.predict(x_test)
```

```
In [170... from sklearn.metrics import confusion_matrix
from sklearn.metrics import accuracy_score
print(confusion_matrix(y_test, y_predict))
```

```
[[ 940  927]
 [ 618 1480]]
```

```
In [171... accuracy_score(y_test, y_predict)*100
```

```
Out[171]: 61.03404791929382
```

Improving the model by changing the kernels

Using the "rbf" kernel

```
In [172... clf = svm.SVC(kernel='rbf', gamma =0.3)
```

```
#clf = svm.LinearSVC(C=1)
clf.fit(x_train, y_train)
y_predict = clf.predict(x_test)
confusion_matrix(y_test, y_predict)
```

```
Out[172]: array([[1084,  783],
               [ 617, 1481]])
```

```
In [173]: accuracy_score(y_test, y_predict)*100
```

```
Out[173]: 64.69104665825978
```

Using the 'poly kernel' kernel

```
In [174]: clf = svm.SVC(kernel='poly', degree=8)
#clf = svm.LinearSVC(C=1)
clf.fit(x_train, y_train)
y_predict = clf.predict(x_test)
print(confusion_matrix(y_test, y_predict))
```

```
[[1217  650]
 [ 754 1344]]
```

```
In [59]: print(accuracy_score(y_test, y_predict)*100)
```

```
64.59016393442623
```

The highly correlated 16 variables were removed from before normalizing the data. The input variables (x variables) were normalized using the MinMaxScaler from sklearn. The data was further split into Training (90%) and Test Data set (10%). The LinearSVC Kernel model showed an accuracy of 61.03%. The accuracy increased to 64.69% when using the 'rbf' kernel with gamma = 0.3. The Accuracy was seen to be 64.59% by changing the Kernel to 'poly' kernel with degrees = 8.

Part III: Dimensionality Reduction using Principal Component Analysis

```
In [175]: import pandas as pd
news = pd.read_csv('OnlineNewsPopularity_for_python.csv')
news.head()
```

```
Out[175]:
```

	url	timedelta	n_tokens_title	n_tokens_content	n_unique_tokens
0	http://mashable.com/2013/01/07/amazon-instant-...	731	12	219	0.663594
1	http://mashable.com/2013/01/07/ap-samsung-spon...	731	9	255	0.604743
2	http://mashable.com/2013/01/07/apple-40-billio...	731	9	211	0.575130
3	http://mashable.com/2013/01/07/astronaut-notre...	731	9	531	0.503788
4	http://mashable.com/2013/01/07/att-u-verse-apps/	731	13	1072	0.415646

5 rows x 61 columns

```
In [176]: news.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

RangeIndex: 39644 entries, 0 to 39643

Data columns (total 61 columns):

#	Column	Non-Null Count	Dtype
---	-----	-----	-----
0	url	39644 non-null	object
1	timedelta	39644 non-null	int64
2	n_tokens_title	39644 non-null	int64
3	n_tokens_content	39644 non-null	int64
4	n_unique_tokens	39644 non-null	float64
5	n_non_stop_words	39644 non-null	float64
6	n_non_stop_unique_tokens	39644 non-null	float64
7	num_hrefs	39644 non-null	int64
8	num_self_hrefs	39644 non-null	int64
9	num_imgs	39644 non-null	int64
10	num_videos	39644 non-null	int64
11	average_token_length	39644 non-null	float64
12	num_keywords	39644 non-null	int64
13	data_channel_is_lifestyle	39644 non-null	int64
14	data_channel_is_entertainment	39644 non-null	int64
15	data_channel_is_bus	39644 non-null	int64
16	data_channel_is_socmed	39644 non-null	int64
17	data_channel_is_tech	39644 non-null	int64
18	data_channel_is_world	39644 non-null	int64
19	kw_min_min	39644 non-null	int64
20	kw_max_min	39644 non-null	float64
21	kw_avg_min	39644 non-null	float64
22	kw_min_max	39644 non-null	int64
23	kw_max_max	39644 non-null	int64
24	kw_avg_max	39644 non-null	float64
25	kw_min_avg	39644 non-null	float64
26	kw_max_avg	39644 non-null	float64
27	kw_avg_avg	39644 non-null	float64
28	self_reference_min_shares	39644 non-null	float64
29	self_reference_max_shares	39644 non-null	float64
30	self_reference_avg_sharess	39644 non-null	float64
31	weekday_is_monday	39644 non-null	int64
32	weekday_is_tuesday	39644 non-null	int64
33	weekday_is_wednesday	39644 non-null	int64
34	weekday_is_thursday	39644 non-null	int64
35	weekday_is_friday	39644 non-null	int64
36	weekday_is_saturday	39644 non-null	int64
37	weekday_is_sunday	39644 non-null	int64
38	is_weekend	39644 non-null	int64
39	LDA_00	39644 non-null	float64
40	LDA_01	39644 non-null	float64
41	LDA_02	39644 non-null	float64
42	LDA_03	39644 non-null	float64
43	LDA_04	39644 non-null	float64
44	global_subjectivity	39644 non-null	float64
45	global_sentiment_polarity	39644 non-null	float64
46	global_rate_positive_words	39644 non-null	float64
47	global_rate_negative_words	39644 non-null	float64
48	rate_positive_words	39644 non-null	float64
49	rate_negative_words	39644 non-null	float64
50	avg_positive_polarity	39644 non-null	float64
51	min_positive_polarity	39644 non-null	float64
52	max_positive_polarity	39644 non-null	float64
53	avg_negative_polarity	39644 non-null	float64
54	min_negative_polarity	39644 non-null	float64
55	max_negative_polarity	39644 non-null	float64
56	title_subjectivity	39644 non-null	float64
57	title_sentiment_polarity	39644 non-null	float64
58	abs_title_subjectivity	39644 non-null	float64
59	abs_title_sentiment_polarity	39644 non-null	float64
60	shares	39644 non-null	int64

```
dtypes: float64(34), int64(26), object(1)
memory usage: 18.5+ MB
```

```
In [177... # handle goal attribute to binary
popular = news.shares >= 1400
unpopular = news.shares < 1400
news.loc[popular, 'shares'] = 1
news.loc[unpopular, 'shares'] = 0
```

```
In [178... target = news["shares"]
y = target
```

Normalizing the data

```
In [179... from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
scaled = scaler.fit_transform(news.drop(columns = ["shares", "url"], axis =1))
scaled
```

```
Out[179]: array([[ 1.75788035,  0.75744723, -0.69521045, ..., -0.97543219,
        -1.81071884,  0.13891975],
       [ 1.75788035, -0.66165665, -0.61879381, ..., -0.26907618,
         0.83774863, -0.68965812],
       [ 1.75788035, -0.66165665, -0.71219192, ..., -0.26907618,
         0.83774863, -0.68965812],
       ...,
       [-1.61808342, -0.18862202, -0.2218518 , ...,  0.24463728,
        -1.56994907, -0.08705603],
       [-1.61808342, -2.08076053,  0.28759248, ..., -0.26907618,
         0.83774863, -0.68965812],
       [-1.61808342, -0.18862202, -0.82681689, ...,  0.67273184,
        -0.92789635,  0.41511238]])
```

Implementing PCA with selecting 30 components from all initial 58 Variables.

```
In [180... from sklearn.decomposition import PCA
pca = PCA(n_components = 30)
PC = pca.fit(scaled)
PC.components_
```

```
Out[180]: array([[ -0.16132739,  0.0581139 , -0.13404156, ..., -0.08046848,
         0.03406751, -0.07174961],
       [ -0.19286504,  0.07672684,  0.09800518, ..., -0.04403546,
        -0.03249026,  0.07886985],
       [ -0.06891273,  0.01090471, -0.13706484, ...,  0.12710808,
        -0.0733436 ,  0.09586703],
       ...,
       [  0.00838535, -0.03247174,  0.10835193, ..., -0.22710656,
         0.11732152,  0.11457629],
       [  0.09999982, -0.0683428 , -0.0681115 , ...,  0.21921158,
        -0.23191139, -0.17044049],
       [ -0.01377077,  0.15246142, -0.19444994, ...,  0.01398225,
         0.58131374,  0.27219535]])
```

Variance in components

```
In [181... PC.explained_variance_ratio_
```

```
Out[181]: array([0.08275377, 0.06985565, 0.06110151, 0.0508775 , 0.0476545 ,
        0.04397503, 0.04309095, 0.03886672, 0.03606727, 0.03525317,
        0.03393574, 0.03195266, 0.02806568, 0.02339546, 0.02313932,
        0.02094132, 0.02082307, 0.02053415, 0.02022003, 0.019877 ,
        0.01914186, 0.01839665, 0.01771758, 0.01543414, 0.01476358,
        0.01378672, 0.01224695, 0.0114735 , 0.01112032, 0.01035154])
```


As we see above the 30 principal components together explain a little over 80 % of the variance. So we will select all thirty Principal Components

```
In [182... new_feats = pca.fit_transform(scaled)
new_feats
```

```
Out[182]: array([[ -1.35785685,  -1.88425763,  -1.97397881, ...,  -0.11094511,
         0.27523182,  -0.27541342],
       [ 0.1963981 ,  -4.30538314,  -1.13590248, ...,  -0.94971978,
         0.57584067,   0.42601707],
       [-3.7311749 ,  -3.49820523,  -1.735055   , ...,   0.39372428,
        -0.3908806 ,   0.49556552],
       ...,
       [ 0.47529804,   2.6445924 ,   0.06452597, ...,  -0.89615935,
         0.70758756,  -1.12919284],
       [ 3.57385922,   0.24496315,  -3.16682215, ...,  -0.84042448,
         0.53751021,   0.10111069],
       [ 0.78548101,   0.61241234,   2.30935024, ...,   0.298666   ,
         1.35605723,  -0.161829   ]])
```

Creating dataframe with the new features as the 30 principal components from above.

```
In [183... new_data = pd.DataFrame(new_feats, columns = ["PCA1", "PCA2", "PCA3", "PC4", "PC5", "PC6
```

```
In [184... new_data.head()
```

```
Out[184]:
```

	PCA1	PCA2	PCA3	PC4	PC5	PC6	PC7	PC8	PCA9
0	-1.357857	-1.884258	-1.973979	-0.000726	0.431248	-3.606503	1.421048	-1.418715	0.462156
1	0.196398	-4.305383	-1.135902	0.082758	-0.351744	-1.565965	3.365590	0.305649	1.896648
2	-3.731175	-3.498205	-1.735055	-0.078569	-0.386974	-0.860576	2.504444	-0.273746	-0.266942
3	-0.670099	-1.695278	-3.436231	-0.106110	1.187241	-2.169244	1.392527	-1.008122	-0.530892
4	-4.169513	-3.718399	-1.441031	-0.131872	0.376862	0.343934	-2.743607	-1.432815	0.532606

5 rows x 30 columns

SVM using "rbf" kernel after dimensionality reduction

```
In [185... clf = svm.SVC(kernel='rbf', gamma =0.3)
#clf = svm.LinearSVC(C=1)
clf.fit(x_train, y_train)
y_predict = clf.predict(x_test)
confusion_matrix(y_test, y_predict)
```

```
Out[185]: array([[1084,   783],
       [  617, 1481]])
```

```
In [186... print(accuracy_score(y_test, y_predict)*100)
```

64.69104665825978

Sunsupervised model: Using Clustering to determine the best value of K using K-means .

```
In [187... from sklearn.cluster import KMeans
import matplotlib.pyplot as plt
```

```


sse = {}
last_sse = 17592402.70373319
for k in range(1,30):
    kmeans = KMeans(n_clusters=k, random_state=12345, n_init = 25).fit(new_data)
    #print(data["clusters"])
    sse[k] = kmeans.inertia_ # Inertia: Sum of distances of samples to their closest clu
    change_per = (last_sse-kmeans.inertia_)/last_sse*100
    print ('At k= ',k,'The percentage of change in SSE is ',change_per,'%')
    last_sse = kmeans.inertia_
plt.figure()
plt.plot(list(sse.keys()), list(sse.values()))
plt.xlabel("Number of cluster")
plt.ylabel("SSE")
plt.show()

```

```

At k= 1 The percentage of change in SSE is 88.07631058253925 %
At k= 2 The percentage of change in SSE is 6.1770128975934755 %
At k= 3 The percentage of change in SSE is 5.352399147817898 %
At k= 4 The percentage of change in SSE is 5.196874067232386 %
At k= 5 The percentage of change in SSE is 5.907706464446229 %
At k= 6 The percentage of change in SSE is 5.588287251806153 %
At k= 7 The percentage of change in SSE is 3.213208487607857 %
At k= 8 The percentage of change in SSE is 4.181117909254028 %
At k= 9 The percentage of change in SSE is 3.8599610578537344 %
At k= 10 The percentage of change in SSE is 2.8623620777938474 %
At k= 11 The percentage of change in SSE is 2.5649047197786743 %
At k= 12 The percentage of change in SSE is 2.802892920114828 %
At k= 13 The percentage of change in SSE is 2.5372429036207804 %
At k= 14 The percentage of change in SSE is 2.255100568286224 %
At k= 15 The percentage of change in SSE is 3.447060012527133 %
At k= 16 The percentage of change in SSE is 0.4836430128399092 %
At k= 17 The percentage of change in SSE is 3.088310387503831 %
At k= 18 The percentage of change in SSE is 1.1019768357863073 %
At k= 19 The percentage of change in SSE is 1.6371358128957423 %
At k= 20 The percentage of change in SSE is 1.7602885621822693 %
At k= 21 The percentage of change in SSE is 0.8808199240166601 %
At k= 22 The percentage of change in SSE is 1.5642818718455027 %
At k= 23 The percentage of change in SSE is 0.37656337146127855 %
At k= 24 The percentage of change in SSE is 0.8745224075705327 %
At k= 25 The percentage of change in SSE is 2.0028257085055117 %
At k= 26 The percentage of change in SSE is 0.6345199003790533 %
At k= 27 The percentage of change in SSE is 0.8788147395915004 %
At k= 28 The percentage of change in SSE is 0.6613765648575097 %
At k= 29 The percentage of change in SSE is 0.6493573618314269 %

```



From the graph it looks like values of k from 5 to 9 can be a good range.

Using KNN classifier which is a Supervised Learning Algorithm

Using k=2, Randomizing the data with seed 12345

```
In [188... import random
target = news["shares"]
random.seed(12345)
indx = random.sample(range(0, 1000), 1000)
new_data_rand = new_data.iloc[indx]
target_rand = target.iloc[indx]
```

Splitting into 80% training data and 20% testing data

```
In [189... from sklearn.model_selection import train_test_split

y = target_rand
x = new_data_rand

x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.20, random_state=1)
```

Using KNN algorithm for Classification with k=2

```
In [190... from sklearn.neighbors import KNeighborsClassifier
neigh = KNeighborsClassifier(n_neighbors=2)
model = neigh.fit(x, y)
y_predict = model.predict(x_test)
```

/Users/pratik/opt/anaconda3/lib/python3.9/site-packages/sklearn/neighbors/_classification.py:228: FutureWarning: Unlike other reduction functions (e.g. `skew`, `kurtosis`), the

default behavior of `mode` typically preserves the axis it acts along. In SciPy 1.11.0, this behavior will change: the default value of `keepdims` will become False, the `axis` over which the statistic is taken will be eliminated, and the value None will no longer be accepted. Set `keepdims` to True or False to avoid this warning.

```
mode, _ = stats.mode(_y[neigh_ind, k], axis=1)
```

```
In [191... from sklearn.metrics import confusion_matrix
from sklearn.metrics import accuracy_score

print(confusion_matrix(y_test, y_predict)*100)

[[9300    0]
 [4100 6600]]
```

```
In [192... print(accuracy_score(y_test, y_predict)*100)

79.5
```

Using K-value of 5

```
In [193... import random
target = news["shares"]
random.seed(12345)
indx = random.sample(range(0, 1000), 1000)
new_data_rand = new_data.iloc[indx]
target_rand = target.iloc[indx]

y = target_rand
x = new_data_rand

x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.20, random_state=1)

neigh = KNeighborsClassifier(n_neighbors=5)
model = neigh.fit(x, y)

y_predict = model.predict(x_test)
```

```
/Users/pratik/opt/anaconda3/lib/python3.9/site-packages/sklearn/neighbors/_classification.py:228: FutureWarning: Unlike other reduction functions (e.g. `skew`, `kurtosis`), the default behavior of `mode` typically preserves the axis it acts along. In SciPy 1.11.0, this behavior will change: the default value of `keepdims` will become False, the `axis` over which the statistic is taken will be eliminated, and the value None will no longer be accepted. Set `keepdims` to True or False to avoid this warning.
mode, _ = stats.mode(_y[neigh_ind, k], axis=1)
```

```
In [194... print(confusion_matrix(y_test, y_predict)*100)
print(accuracy_score(y_test, y_predict)*100)

[[7200 2100]
 [2800 7900]]
75.5
```

Using k-value of 8

```
In [195... import random
target = news["shares"]
random.seed(12345)
indx = random.sample(range(0, 1000), 1000)
new_data_rand = new_data.iloc[indx]
target_rand = target.iloc[indx]
```

```

y = target_rand
x = new_data_rand

x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.20, random_state=1)

neigh = KNeighborsClassifier(n_neighbors=8)
model = neigh.fit(x, y)

y_predict = model.predict(x_test)

print(confusion_matrix(y_test, y_predict)*100)
print(accuracy_score(y_test, y_predict)*100)

[[7600 1700]
 [4700 6000]]
68.0

```

/Users/pratik/opt/anaconda3/lib/python3.9/site-packages/sklearn/neighbors/_classification.py:228: FutureWarning: Unlike other reduction functions (e.g. `skew`, `kurtosis`), the default behavior of `mode` typically preserves the axis it acts along. In SciPy 1.11.0, this behavior will change: the default value of `keepdims` will become False, the `axis` over which the statistic is taken will be eliminated, and the value None will no longer be accepted. Set `keepdims` to True or False to avoid this warning.

```

mode, _ = stats.mode(_y[neigh_ind, k], axis=1)

```

Using K-20

In [196...

```

import random
target = news["shares"]
random.seed(12345)
indx = random.sample(range(0, 1000), 1000)
new_data_rand = new_data.iloc[indx]
target_rand = target.iloc[indx]

y = target_rand
x = new_data_rand

x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.20, random_state=1)

neigh = KNeighborsClassifier(n_neighbors=21)
model = neigh.fit(x, y)

y_predict = model.predict(x_test)

print(confusion_matrix(y_test, y_predict)*100)
print(accuracy_score(y_test, y_predict)*100)

[[6400 2900]
 [4200 6500]]
64.5

```

/Users/pratik/opt/anaconda3/lib/python3.9/site-packages/sklearn/neighbors/_classification.py:228: FutureWarning: Unlike other reduction functions (e.g. `skew`, `kurtosis`), the default behavior of `mode` typically preserves the axis it acts along. In SciPy 1.11.0, this behavior will change: the default value of `keepdims` will become False, the `axis` over which the statistic is taken will be eliminated, and the value None will no longer be accepted. Set `keepdims` to True or False to avoid this warning.

```

mode, _ = stats.mode(_y[neigh_ind, k], axis=1)

```

As seen above the accuracy for k=2 is 80%, K=5 is 75.5%, k=8 is 69.5% and k=21 is 63.5%. However, from the graph generated above, it looks like k=8 is a better value as the sse is small compared to K=2 and k=5. A slight elbow can be seen from the graph

as seen from $k=5$ to $k=8$ with a good percentage change in sse. The percentage in sse after $k=9$ is very less comparatively. I would like to use value of $K=8$.

Conclusion:

Amongst all the algorithms used, the KNN Classifier algorithm after Dimensionality Reduction gave us the maximum accuracy for predicting the Popular news articles.